

Importing libraries

```
!pip install tensorflow
```

```
Requirement already satisfied: tensorflow in
./env/lib/python3.10/site-packages (2.13.0rc1)
Requirement already satisfied: tensorflow-macos==2.13.0-rc1 in
./env/lib/python3.10/site-packages (from tensorflow) (2.13.0rc1)
Requirement already satisfied: termcolor>=1.1.0 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (2.3.0)
Requirement already satisfied: packaging in ./env/lib/python3.10/site-
packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (22.0)
Requirement already satisfied: setuptools in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (65.6.3)
Requirement already satisfied: six>=1.12.0 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (1.16.0)
Requirement already satisfied: flatbuffers>=23.1.21 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (23.5.26)
Requirement already satisfied: h5py>=2.9.0 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (3.8.0)
Requirement already satisfied: wrapt>=1.11.0 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (1.15.0)
Requirement already satisfied: gast<=0.4.0,>=0.2.1 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (0.4.0)
Requirement already satisfied: numpy>=1.22 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (1.23.5)
Requirement already satisfied: typing-extensions>=3.6.6 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (4.4.0)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!
=4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in
./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1-
>tensorflow) (4.23.3)
Requirement already satisfied: opt-einsum>=2.3.2 in
```

./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (3.3.0)
Requirement already satisfied: tensorflow-estimator<2.14,>=2.13.0rc0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (2.13.0rc0)
Requirement already satisfied: tensorboard<2.14,>=2.13 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (2.13.0)
Requirement already satisfied: google-pasta>=0.1.1 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (16.0.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (1.54.2)
Requirement already satisfied: keras<2.14,>=2.13.1rc0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (2.13.1rc0)
Requirement already satisfied: astunparse>=1.6.0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (1.6.3)
Requirement already satisfied: absl-py>=1.0.0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (1.4.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (0.38.4)
Requirement already satisfied: markdown>=2.6.8 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (3.4.3)
Requirement already satisfied: requests<3,>=2.21.0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (2.28.1)
Requirement already satisfied: werkzeug>=1.0.1 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (2.3.6)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (0.7.1)
Requirement already satisfied: google-auth-oauthlib<1.1,>=0.5 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (1.0.0)
Requirement already satisfied: google-auth<3,>=1.6.3 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (2.20.0)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in ./env/lib/python3.10/site-packages (from tensorflow-macos==2.13.0-rc1->tensorflow) (5.3.0)

```

(5.3.1)
Requirement already satisfied: urllib3<2.0 in
./env/lib/python3.10/site-packages (from google-auth<3,>=1.6.3-
>tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-rc1->tensorflow)
(1.26.14)
Requirement already satisfied: pyasn1-modules>=0.2.1 in
./env/lib/python3.10/site-packages (from google-auth<3,>=1.6.3-
>tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-rc1->tensorflow)
(0.3.0)
Requirement already satisfied: rsa<5,>=3.1.4 in
./env/lib/python3.10/site-packages (from google-auth<3,>=1.6.3-
>tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-rc1->tensorflow)
(4.9)
Requirement already satisfied: requests-oauthlib>=0.7.0 in
./env/lib/python3.10/site-packages (from google-auth-
oauthlib<1.1,>=0.5->tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-
rc1->tensorflow) (1.3.1)
Requirement already satisfied: certifi>=2017.4.17 in
./env/lib/python3.10/site-packages (from requests<3,>=2.21.0-
>tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-rc1->tensorflow)
(2022.6.15)
Requirement already satisfied: charset-normalizer<3,>=2 in
./env/lib/python3.10/site-packages (from requests<3,>=2.21.0-
>tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-rc1->tensorflow)
(2.0.4)
Requirement already satisfied: idna<4,>=2.5 in
./env/lib/python3.10/site-packages (from requests<3,>=2.21.0-
>tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-rc1->tensorflow)
(3.4)
Requirement already satisfied: MarkupSafe>=2.1.1 in
./env/lib/python3.10/site-packages (from werkzeug>=1.0.1-
>tensorboard<2.14,>=2.13->tensorflow-macos==2.13.0-rc1->tensorflow)
(2.1.1)
Requirement already satisfied: pyasn1<0.6.0,>=0.4.6 in
./env/lib/python3.10/site-packages (from pyasn1-modules>=0.2.1-
>google-auth<3,>=1.6.3->tensorboard<2.14,>=2.13->tensorflow-
macos==2.13.0-rc1->tensorflow) (0.5.0)
Requirement already satisfied: oauthlib>=3.0.0 in
./env/lib/python3.10/site-packages (from requests-oauthlib>=0.7.0-
>google-auth-oauthlib<1.1,>=0.5->tensorboard<2.14,>=2.13->tensorflow-
macos==2.13.0-rc1->tensorflow) (3.2.2)

```

```

import pandas as pd
import numpy as np
from scipy import stats
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

```

```

from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score
from sklearn.neural_network import MLPRegressor
import tensorflow as tf
import time
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error

```

importing dataset

```

df = pd.read_excel("Dataset.xlsx", names=[ "Time_Interval", "NO",
"NO2", "SO2", "Air_temp", "Rel_Humidity", "Wind_Dir", "Wind_speed",
"Lag1", "Lag2", "PM"],header=None);

```

df

| | | Time_Interval | NO | NO2 | SO2 | Air_temp | |
|----------------|------------|---------------|------|-------|-----|----------|-------|
| Rel_Humidity \ | | | | | | | |
| 0 | 2020-01-01 | 00:00:00 | 5.25 | 14.80 | NaN | 18.0 | 76.00 |
| 1 | 2020-01-01 | 01:00:00 | 4.90 | 15.05 | NaN | 18.0 | 77.60 |
| 2 | 2020-01-01 | 02:00:00 | 2.35 | 14.10 | 0.0 | 18.0 | 79.10 |
| 3 | 2020-01-01 | 03:00:00 | 1.55 | 16.00 | 0.0 | 18.0 | 80.20 |
| 4 | 2020-01-01 | 04:00:00 | 1.60 | 16.75 | 0.0 | 18.0 | 80.35 |
| ... | | ... | ... | ... | ... | ... | ... |
| 26300 | 2022-12-31 | 20:00:00 | 0.50 | 2.10 | 0.2 | NaN | NaN |
| 26301 | 2022-12-31 | 21:00:00 | 0.50 | 1.35 | 0.2 | NaN | NaN |
| 26302 | 2022-12-31 | 22:00:00 | 0.50 | 1.20 | 0.4 | NaN | NaN |
| 26303 | | NaN | NaN | NaN | NaN | NaN | NaN |
| 26304 | | NaN | NaN | NaN | NaN | NaN | NaN |

| | Wind_Dir | Wind_speed | Lag1 | Lag2 | PM |
|---|----------|------------|-------|-------|-------|
| 0 | 223.5 | 2.20 | NaN | NaN | 10.60 |
| 1 | 217.0 | 2.10 | 10.60 | NaN | 8.60 |
| 2 | 212.0 | 1.75 | 8.60 | 10.60 | 7.85 |
| 3 | 188.0 | 1.20 | 7.85 | 8.60 | 8.60 |
| 4 | 172.5 | 1.05 | 8.60 | 7.85 | 9.30 |

```

...      ...      ...      ...      ...
26300      NaN      NaN      -0.25      -0.75      0.90
26301      NaN      NaN      0.90      -0.25      2.00
26302      NaN      NaN      2.00      0.90      1.95
26303      NaN      NaN      1.95      2.00      NaN
26304      NaN      NaN      NaN      1.95      NaN

```

```
[26305 rows x 11 columns]
```

```
df['Time_Interval'] = pd.to_datetime(df['Time_Interval']) # Convert
the timestamp column to datetime format
```

```
df.info()
```

```


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26305 entries, 0 to 26304
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Time_Interval          26303 non-null  datetime64[ns]
1   NO                      22790 non-null  float64
2   NO2                     22948 non-null  float64
3   SO2                     23766 non-null  float64
4   Air_temp               24776 non-null  float64
5   Rel_Humidity           24912 non-null  float64
6   Wind_Dir               23653 non-null  float64
7   Wind_speed             24727 non-null  float64
8   Lag1                   22798 non-null  float64
9   Lag2                   22798 non-null  float64
10  PM                      22798 non-null  float64
dtypes: datetime64[ns](1), float64(10)
memory usage: 2.2 MB

```

```

plt.plot(df["Time_Interval"], df["PM"])
plt.xlabel("Time")
plt.ylabel("PM 2.5")
plt.title("PM 2.5 Trend Over Time")
plt.xticks(rotation=45)
plt.show()

```



```
# Identify and handle missing data
```

```
missing_data = df.isnull().sum()
```

```
missing_data
```

```
Time_Interval      2
```

```
N0                 3515
```

```
N02                3357
```

```
S02                2539
```

```
Air_temp           1529
```

```
Rel_Humidity        1393
```

```
Wind_Dir            2652
```

```
Wind_speed           1578
```

```
Lag1                3507
```

```
Lag2                3507
```

```
PM                  3507
```

```
dtype: int64
```

```
# Replace negative values with NaN in PM2.5 column as mentioned by  
SARA zandi in her presentation
```

```
df['PM'] = np.where(df['PM'] < 0, np.nan, df['PM'])
```

```
missing_data = df.isnull().sum()
missing_data
```

```
Time_Interval      2
N0                  3515
N02                 3357
S02                 2539
Air_temp            1529
Rel_Humidity        1393
Wind_Dir            2652
Wind_speed          1578
Lag1                3507
Lag2                3507
PM                  4722
dtype: int64
```

```
# Replace missing values with the mean, median and mode imputation
```

```
mean_df = df.fillna(df.mean(), inplace=False)
median_df=df.fillna(df.median(), inplace=False)
mode_df= df.fillna(df.mode().iloc[0])
```

```
/var/folders/r9/ftlfg09n2rbgf2q9jy3b7j7h0000gn/T/
ipykernel_21959/2800757332.py:2: FutureWarning: DataFrame.mean and
DataFrame.median with numeric_only=None will include datetime64 and
datetime64tz columns in a future version.
```

```
mean_df = df.fillna(df.mean(), inplace=False)
/var/folders/r9/ftlfg09n2rbgf2q9jy3b7j7h0000gn/T/ipykernel_21959/28007
57332.py:3: FutureWarning: DataFrame.mean and DataFrame.median with
numeric_only=None will include datetime64 and datetime64tz columns in
a future version.
```

```
median_df=df.fillna(df.median(), inplace=False)
```

```
### plotting boxplot for all three imputation method to compare and see
which works best
```

```
#ploting mean impute
```


```
fig, ax = plt.subplots(figsize=(10, 6))
mean_df.boxplot(ax=ax)
plt.title("Data Filled with mean impute")
# # Set the labels for the x-axis and y-axis
plt.xlabel("Columns")
plt.ylabel("Values")
# # Display the plot
plt.show()
```

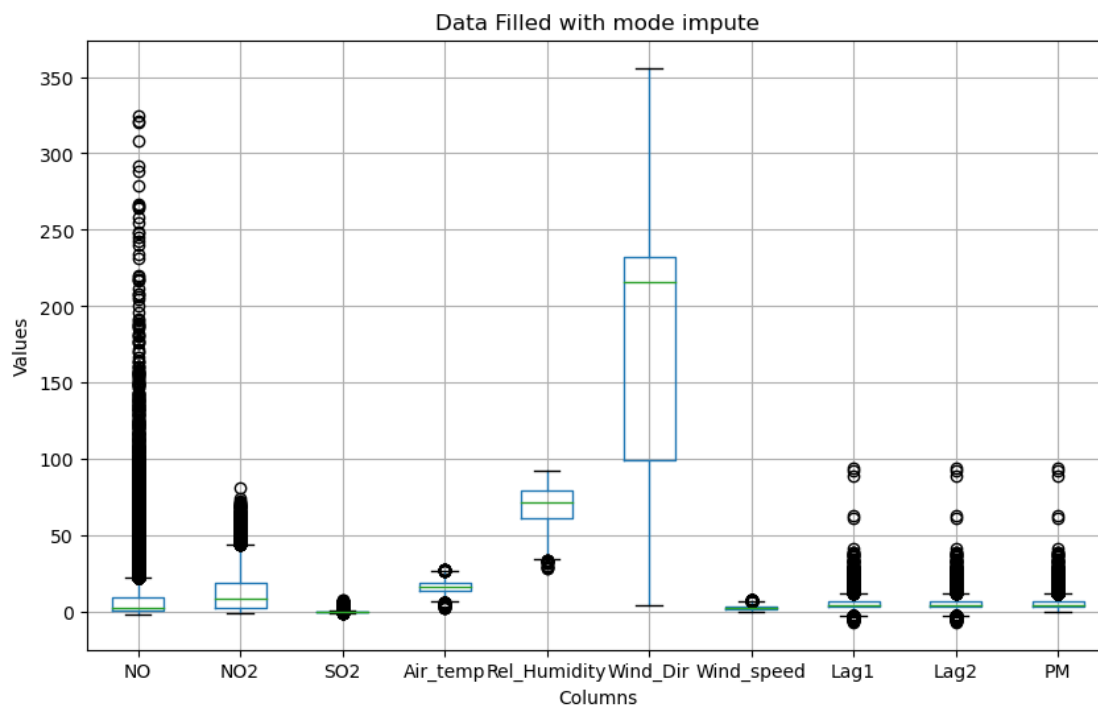

```
#ploting median impute
```

```
fig, ax = plt.subplots(figsize=(10, 6))
median_df.boxplot(ax=ax)
plt.title("Data Filled with median impute")
# # Set the labels for the x-axis and y-axis
plt.xlabel("Columns")
```

```
plt.ylabel("Values")
# # Display the plot
plt.show()
```

```
#ploting mode impute
fig, ax = plt.subplots(figsize=(10, 6))
mode_df.boxplot(ax=ax)
plt.title("Data Filled with mode impute")
# # Set the labels for the x-axis and y-axis
plt.xlabel("Columns")
plt.ylabel("Values")
# # Display the plot
plt.show()
```





#Since there is not significant difference when comparing we'll go with mean imputation

Identify and handle outliers

```
outliers = mean_df[["NO", "NO2", "SO2", "Air_temp", "Rel_Humidity",
"Wind_Dir", "Wind_speed", "Lag1", "Lag2", "PM"]].apply(lambda x:
```

```
np.abs(x - x.mean()) / x.std() > 3).sum()
outliers
```

```
NO          539
N02         361
S02         493
Air_temp    45
Rel_Humidity 14
Wind_Dir     0
Wind_speed  55
Lag1        351
Lag2        351
PM          390
dtype: int64
```

```
# Define the columns with outliers
```

```
outlier_columns = [ 'NO', 'N02', 'S02', 'Air_temp', 'Rel_Humidity',
'Wind_speed', 'PM', 'Lag1', 'Lag2']
```

```
# Calculate z-scores for outlier detection and removal
```

```
z_scores = np.abs(stats.zscore(mean_df[outlier_columns]))
threshold = 3
```

```
# Remove rows with outliers
```

```
mean_df = mean_df[(z_scores < threshold).all(axis=1)]
```

```
# Verify the updated dataset
```

```
(mean_df.shape)
```

```
(24549, 11)
```


```
mean_df
```

| | Time_Interval | NO | N02 | S02 | Air_temp \ |
|-----------|---------------------|------------|------------|-----------|------------|
| 0 | 2020-01-01 00:00:00 | 5.250000 | 14.800000 | 0.326079 | 18.000000 |
| 1 | 2020-01-01 01:00:00 | 4.900000 | 15.050000 | 0.326079 | 18.000000 |
| 2 | 2020-01-01 02:00:00 | 2.350000 | 14.100000 | 0.000000 | 18.000000 |
| 3 | 2020-01-01 03:00:00 | 1.550000 | 16.000000 | 0.000000 | 18.000000 |
| 4 | 2020-01-01 04:00:00 | 1.600000 | 16.750000 | 0.000000 | 18.000000 |
| ... | ... | ... | ... | ... | ... |
| 26300 | 2022-12-31 20:00:00 | 0.500000 | 2.100000 | 0.200000 | 16.567666 |
| 26301 | 2022-12-31 21:00:00 | 0.500000 | 1.350000 | 0.200000 | 16.567666 |
| 26302 | 2022-12-31 22:00:00 | 0.500000 | 1.200000 | 0.400000 | 16.567666 |
| 26303 | NaT | 9.995042 | 14.425854 | 0.326079 | 16.567666 |
| 26304 | NaT | 9.995042 | 14.425854 | 0.326079 | 16.567666 |
| | | | | | |
| | Rel_Humidity | Wind_Dir | Wind_speed | Lag1 | Lag2 |
| PM | | | | | |
| 0 | 76.000000 | 223.500000 | 2.200000 | 5.428717 | 5.428717 |
| 10.600000 | | | | | |
| 1 | 77.600000 | 217.000000 | 2.100000 | 10.600000 | 5.428717 |

| | | | | | |
|---------|-----------|------------|----------|-----------|-----------|
| 8.60000 | | | | | |
| 2 | 79.100000 | 212.000000 | 1.750000 | 8.600000 | 10.600000 |
| 7.85000 | | | | | |
| 3 | 80.200000 | 188.000000 | 1.200000 | 7.850000 | 8.600000 |
| 8.60000 | | | | | |
| 4 | 80.350000 | 172.500000 | 1.050000 | 8.600000 | 7.850000 |
| 9.30000 | | | | | |
| ... | ... | ... | ... | ... | ... |
| ... | | | | | |
| 26300 | 69.785531 | 175.673297 | 2.777638 | -0.250000 | -0.750000 |
| 0.90000 | | | | | |
| 26301 | 69.785531 | 175.673297 | 2.777638 | 0.900000 | -0.250000 |
| 2.00000 | | | | | |
| 26302 | 69.785531 | 175.673297 | 2.777638 | 2.000000 | 0.900000 |
| 1.95000 | | | | | |
| 26303 | 69.785531 | 175.673297 | 2.777638 | 1.950000 | 2.000000 |
| 5.80859 | | | | | |
| 26304 | 69.785531 | 175.673297 | 2.777638 | 5.428717 | 1.950000 |
| 5.80859 | | | | | |

[24549 rows x 11 columns]

```
plt.plot(mean_df["Time_Interval"], mean_df["PM"])
plt.xlabel("Time")
plt.ylabel("PM 2.5")
plt.title("PM 2.5 Trend Over Time after data cleaning")
plt.xticks(rotation=45)
plt.show()
```



```


correlation_matrix = mean_df.corr()
target_correlation = correlation_matrix['PM'].drop('PM')
highest_correlation_predictor = target_correlation.idxmax()

/var/folders/r9/ftlfg09n2rbgf2q9jy3b7j7h0000gn/T/
ipykernel_21959/1486501280.py:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it
will default to False. Select only valid columns or specify the value
of numeric_only to silence this warning.
    correlation_matrix = mean_df.corr()

# Plotting the highest correlated predictor against the target
variable
sns.scatterplot(x=highest_correlation_predictor, y='PM', data=mean_df)

plt.xlabel(highest_correlation_predictor)
plt.ylabel('Target')
plt.title(f'Highest Correlation: {highest_correlation_predictor} vs
Target')
plt.gcf().set_size_inches((20, 10))
plt.show()

```



```
mean_df.describe()    #Statistical summary
```

| | NO | NO2 | S02 | Air_temp |
|----------------|--------------|--------------|--------------|--------------|
| Rel_Humidity \ | | | | |
| count | 24549.000000 | 24549.000000 | 24549.000000 | 24549.000000 |
| mean | 7.415754 | 13.285116 | 0.254007 | 16.727087 |
| std | 9.047286 | 10.239914 | 0.261966 | 3.630268 |
| min | -1.750000 | -0.900000 | -1.000000 | 5.500000 |
| 25% | 1.100000 | 4.900000 | 0.050000 | 14.000000 |
| 50% | 4.550000 | 12.500000 | 0.250000 | 16.567666 |
| 75% | 9.995042 | 17.600000 | 0.350000 | 19.000000 |
| max | 64.700000 | 49.500000 | 1.800000 | 27.500000 |

| | Wind_Dir | Wind_speed | Lag1 | Lag2 |
|-------|--------------|--------------|--------------|--------------|
| PM | | | | |
| count | 24549.000000 | 24549.000000 | 24549.000000 | 24549.000000 |
| mean | 174.269569 | 2.824292 | 5.004147 | 5.062617 |
| std | 87.084853 | 1.456024 | 3.243318 | 3.315584 |
| min | 4.000000 | 0.250000 | -6.550000 | -6.550000 |
| 25% | 92.500000 | 1.700000 | 2.950000 | 2.950000 |

| | | | | |
|-----------|------------|----------|-----------|-----------|
| 3.400000 | | | | |
| 50% | 197.500000 | 2.750000 | 5.250000 | 5.300000 |
| 5.750000 | | | | |
| 75% | 232.500000 | 3.750000 | 6.500000 | 6.600000 |
| 6.500000 | | | | |
| max | 356.000000 | 7.150000 | 17.500000 | 17.550000 |
| 17.100000 | | | | |

Compute the Pearson correlation matrix

```
correlation_matrix = mean_df.corr()
```

Plot the correlation matrix

```
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
```


```
plt.title('Pearson Correlation Matrix')
```

```
plt.show()
```

/var/folders/r9/ftlfg09n2rbgf2q9jy3b7j7h0000gn/T/

ipykernel_21959/3291743950.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_matrix = mean_df.corr()
```



Based on the correlation coefficients: It appears that Lag1 and Lag2 have the highest positive correlations with PM2.5. Other variables such as NO, NO2, and SO2 also show moderate positive correlations. On the other hand, variables like Air_temp, Rel_Humidity and Wind_Dir have weak negative correlations with PM2.5.

Selecting the features with the highest absolute correlation coefficients

```
selected_features = ['NO', 'NO2', 'SO2', 'Lag1', 'Lag2']
```

Create a new DataFrame with the selected features

```
df_selected = mean_df[selected_features + ['PM']]
```

Split the dataset into features (X) and target variable (y)

```
X = df_selected.drop('PM', axis=1)
```

```
y = df_selected['PM']
```

Split the data into 70% training and 30% testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
```

```

reg_model = LinearRegression()
reg_model.fit(X_train, y_train)
y_pred = reg_model.predict(X_test)

# Calculate evaluation metrics
rmse_reg = mean_squared_error(y_test, y_pred, squared=False)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print the evaluation metrics
print("Linear Regression Model Performance:")
print("RMSE:", rmse_reg)
print("MAE:", mae)
print("R2 Score:", r2)

Linear Regression Model Performance:
RMSE: 2.1659689714490153
MAE: 1.6162891175679295
R2 Score: 0.44726052657038085

import statsmodels.formula.api as smf
mod1 = smf.ols("PM ~ NO + NO2 + SO2 + Lag1 + Lag2", data=mean_df)
mod1_res = mod1.fit()
mod1_res.summary()

```

```

<class 'statsmodels.iolib.summary.Summary'>
"""

```

OLS Regression Results

```

=====
=====
Dep. Variable:          PM    R-squared:
0.449
Model:                OLS    Adj. R-squared:
0.449
Method:             Least Squares    F-statistic:
3999.
Date:              Fri, 16 Jun 2023    Prob (F-statistic):
0.00
Time:              20:54:42    Log-Likelihood:
-53683.
No. Observations:    24549    AIC:
1.074e+05
Df Residuals:        24543    BIC:
1.074e+05
Df Model:            5

Covariance Type:      nonrobust

=====
=====

```


| | coef | std err | t | P> t | [0.025 |
|----------------|---------|----------|-------------------|-------|--------|
| 0.975] | | | | | |
| ----- | | | | | |
| ----- | | | | | |
| Intercept | 2.5681 | 0.032 | 81.078 | 0.000 | 2.506 |
| 2.630 | | | | | |
| N0 | 0.0234 | 0.002 | 10.386 | 0.000 | 0.019 |
| 0.028 | | | | | |
| N02 | 0.0137 | 0.002 | 7.197 | 0.000 | 0.010 |
| 0.017 | | | | | |
| S02 | 0.1963 | 0.059 | 3.343 | 0.001 | 0.081 |
| 0.311 | | | | | |
| Lag1 | 0.7165 | 0.006 | 116.403 | 0.000 | 0.704 |
| 0.729 | | | | | |
| Lag2 | -0.2289 | 0.006 | -38.293 | 0.000 | -0.241 |
| -0.217 | | | | | |
| ===== | | | | | |
| ===== | | | | | |
| Omnibus: | | 2021.777 | Durbin-Watson: | | |
| 1.600 | | | | | |
| Prob(Omnibus): | | 0.000 | Jarque-Bera (JB): | | |
| 3153.086 | | | | | |
| Skew: | | 0.638 | Prob(JB): | | |
| 0.00 | | | | | |
| Kurtosis: | | 4.205 | Cond. No. | | |
| 89.3 | | | | | |
| ===== | | | | | |
| ===== | | | | | |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 """

Multilayer Perceptron (MLP)

```
# Part 2: Single hidden layer with k = 25 neurons
# Create an MLPRegressor with default values and a single hidden layer
of 25 neurons
mlp = MLPRegressor(hidden_layer_sizes=(25, ), random_state=42)

# Train the model using the training dataset
mlp.fit(X_train, y_train)

MLPRegressor(hidden_layer_sizes=(25, ), random_state=42)

# Make predictions on the testing dataset
y_pred_mlp = mlp.predict(X_test)
```

```

# Evaluate the performance using mean squared error (MSE)
mse = mean_squared_error(y_test, y_pred_mlp)
print("Mean Squared Error with default learning rate:", mse)

Mean Squared Error with default learning rate: 4.2905703603859955

mse = mean_squared_error(y_test, y_pred_mlp)
rmse_mlp = mean_squared_error(y_test, y_pred_mlp, squared=False)
mae = mean_absolute_error(y_test, y_pred_mlp)
r2 = r2_score(y_test, y_pred_mlp)

# Print the evaluation metrics
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse_mlp)
print("Mean Absolute Error:", mae)
print("R2 Score:", r2)

Mean Squared Error: 4.2905703603859955
Root Mean Squared Error: 2.0713691994393457
Mean Absolute Error: 1.559268079477574
R2 Score: 0.49448849168583275

k = 25 # Total number of neurons across two hidden layers
# Part 3: Experimenting with two hidden layers

best_mse = float('inf')
best_hidden_layers = None

# Iterate over different configurations of neurons in two hidden
layers
for neurons_in_first_layer in range(k-1, 0, -1):
    neurons_in_second_layer = k - neurons_in_first_layer

    # Create an MLPRegressor with two hidden layers
    mlp = MLPRegressor(hidden_layer_sizes=(neurons_in_first_layer,
neurons_in_second_layer), random_state=42)

    # Train the model using the training dataset
    mlp.fit(X_train, y_train)

    # Make predictions on the testing dataset
    y_pred_mlp = mlp.predict(X_test)

    # Evaluate the performance using mean squared error (MSE)
    mse = mean_squared_error(y_test, y_pred_mlp)

    # Check if this configuration gives the lowest MSE
    if mse < best_mse:
        best_mse = mse
        best_hidden_layers = (neurons_in_first_layer,
neurons_in_second_layer)

```

```
print("Best hidden layer configuration:", best_hidden_layers)
print("Lowest Mean Squared Error:", best_mse)
```

Best hidden layer configuration: (15, 10)
Lowest Mean Squared Error: 4.258910014126226

*###the best configuration consists of two hidden layers with 15 neurons in the first layer and 10 neurons in the second layer.
##The lowest MSE achieved with this configuration is approximately 4.259.*

##we can conclude that the architecture with (15, 10) neurons in the two hidden layers performs slightly better than the default architecture in part .

Long Short-Term Memory (LSTM)

Normalize the data

```
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Reshape the input data for LSTM

```
X_train_reshaped = X_train_scaled.reshape(X_train_scaled.shape[0],
X_train_scaled.shape[1], 1)
X_test_reshaped = X_test_scaled.reshape(X_test_scaled.shape[0],
X_test_scaled.shape[1], 1)
```

##apply LSTM model

```
model = Sequential()
model.add(LSTM(25, activation='relu',
input_shape=(X_train_reshaped.shape[1], 1)))
model.add(Dense(1))
```

```
learning_rate = 0.01
batch_size = 4
```

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
              loss='mean_squared_error')
```


```
history = model.fit(X_train_reshaped, y_train,
validation_data=(X_test_reshaped, y_test), batch_size=batch_size,
epochs=32, verbose=0)
```

WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.

WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.

```
train_loss = history.history['loss']  
test_loss = model.evaluate(X_test_resaped, y_test, verbose=0)  
  
plt.plot(train_loss, label='Train')  
plt.axhline(y=test_loss, color='r', linestyle='--', label='Test')  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```

```
best_epoch = np.argmin(test_loss) + 1
```



```
summary_stats = {  
    'Mean': np.mean(train_loss),  
    'Standard Deviation': np.std(train_loss),  
    'Minimum': np.min(train_loss),  
}
```

```

        'Best Epoch': best_epoch,
        'Maximum': np.max(train_loss),
        'Runtime': history.epoch[-1] + 1 # Epochs start from 0, so adding
1 to get the runtime
    }

```

```
summary_stats
```

```

{'Mean': 4.376399333660419,
 'Standard Deviation': 0.21613654042756908,
 'Minimum': 4.248939514160156,
 'Best Epoch': 25,
 'Maximum': 5.378185272216797,
 'Runtime': 26}

```

```
best_epoch = 25
```

```

batch_sizes = [2, 4, 8, 16, 32, 64, 90 , 150]
summary_stats_batch = []

```

```

for batch_size in batch_sizes:
    model = Sequential()
    model.add(LSTM(25, activation='relu',
input_shape=(X_train_resaped.shape[1], 1)))
    model.add(Dense(1))

```

```

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
               loss='mean_squared_error')

```

```

    start_time = time.time()
    history = model.fit(X_train_resaped, y_train,
batch_size=batch_size, epochs=best_epoch, verbose=0)
    end_time = time.time()

```

```

    train_loss = history.history['loss']
    runtime = end_time - start_time

```

```

# Calculate summary statistics
mean_loss = np.mean(train_loss)
std_loss = np.std(train_loss)
min_loss = np.min(train_loss)
max_loss = np.max(train_loss)

```

```

summary_stats_batch.append({
    'Batch Size': batch_size,
    'Mean': mean_loss,

```

```
    'Standard Deviation': std_loss,  
    'Minimum': min_loss,  
    'Maximum': max_loss,  
    'Runtime': runtime  
  })
```

summary_stats_batch

```
WARNING:absl:At this time, the v2.11+ optimizer  
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the  
legacy Keras optimizer instead, located at  
`tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:There is a known slowdown when using v2.11+ Keras  
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,  
i.e., `tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:At this time, the v2.11+ optimizer  
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the  
legacy Keras optimizer instead, located at  
`tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:There is a known slowdown when using v2.11+ Keras  
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,  
i.e., `tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:At this time, the v2.11+ optimizer  
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the  
legacy Keras optimizer instead, located at  
`tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:There is a known slowdown when using v2.11+ Keras  
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,  
i.e., `tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:At this time, the v2.11+ optimizer  
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the  
legacy Keras optimizer instead, located at  
`tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:There is a known slowdown when using v2.11+ Keras  
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,  
i.e., `tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:At this time, the v2.11+ optimizer  
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the  
legacy Keras optimizer instead, located at  
`tf.keras.optimizers.legacy.Adam`.  
WARNING:absl:There is a known slowdown when using v2.11+ Keras  
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,  
i.e., `tf.keras.optimizers.legacy.Adam`.
```

WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.
WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.
WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.
WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.

```
[{'Batch Size': 2,  
  'Mean': 4.4142498016357425,  
  'Standard Deviation': 0.3453042507018569,  
  'Minimum': 4.253872394561768,  
  'Maximum': 6.041283130645752,  
  'Runtime': 166.30285096168518},  
{ 'Batch Size': 4,  
  'Mean': 4.382492961883545,  
  'Standard Deviation': 0.3646838634618605,  
  'Minimum': 4.240434646606445,  
  'Maximum': 6.108152866363525,  
  'Runtime': 81.73005795478821},  
{ 'Batch Size': 8,  
  'Mean': 4.348611583709717,  
  'Standard Deviation': 0.41452505050741245,  
  'Minimum': 4.189599990844727,  
  'Maximum': 6.330445289611816,  
  'Runtime': 40.6296010017395},  
{ 'Batch Size': 16,  
  'Mean': 4.409920616149902,  
  'Standard Deviation': 0.5709795356460061,  
  'Minimum': 4.187167167663574,  
  'Maximum': 7.124744415283203,  
  'Runtime': 21.053802967071533},  
{ 'Batch Size': 32,  
  'Mean': 4.430623512268067,  
  'Standard Deviation': 0.5748771832378529,  
  'Minimum': 4.219583034515381,  
  'Maximum': 7.179099082946777,  
  'Runtime': 11.376978158950806},  
{ 'Batch Size': 64,  
  'Mean': 4.7081369781494145,  
  'Standard Deviation': 0.9844180699680888,  
  'Minimum': 4.225187301635742,  
  'Maximum': 8.683056831359863,
```

```

    'Runtime': 9.567314863204956},
    {'Batch Size': 90,
     'Mean': 4.731443424224853,
     'Standard Deviation': 1.1338990095597246,
     'Minimum': 4.25293493270874,
     'Maximum': 9.498494148254395,
     'Runtime': 8.35860300064087},
    {'Batch Size': 150,
     'Mean': 5.022337608337402,
     'Standard Deviation': 1.5644735628215338,
     'Minimum': 4.248717784881592,
     'Maximum': 11.382460594177246,
     'Runtime': 5.691593885421753}]

```

The batch size of 8 shows a relatively low mean value (4.35), indicating good overall performance in terms of minimizing the cost function. Additionally, it has a moderate standard deviation (0.41), suggesting reasonably consistent results across the runs. Moreover, the runtime for a batch size of 8 is shorter compared to smaller batch sizes (2, 4), making it a reasonable choice in terms of computational efficiency.

```

neuron_counts = [10, 25, 50, 75, 100] # Define the different numbers
of neurons

```

```

results = [] # Store the results for each run

```

```

for neurons in neuron_counts:
    model = Sequential()
    model.add(LSTM(neurons, activation='relu',
input_shape=(X_train_resaped.shape[1], 1)))
    model.add(Dense(1))

```

```

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.01),
              loss='mean_squared_error')

```

```

history = model.fit(X_train_resaped, y_train, batch_size=8,
epochs=26, verbose=0)

```

```

train_loss = history.history['loss']
runtime = len(train_loss) # Runtime is the number of epochs

```

```

result = {
    'Neuron Count': neurons,
    'Mean': np.mean(train_loss),
    'Standard Deviation': np.std(train_loss),
    'Minimum': np.min(train_loss),

```



```

        'Maximum': np.max(train_loss),
        'Runtime': runtime
    }

```

```

    results.append(result)

```

```

WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.
WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.
WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.
WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.
WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.
WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.
WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.
WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.
WARNING:absl:At this time, the v2.11+ optimizer
`tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the
legacy Keras optimizer instead, located at
`tf.keras.optimizers.legacy.Adam`.
WARNING:absl:There is a known slowdown when using v2.11+ Keras
optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer,
i.e., `tf.keras.optimizers.legacy.Adam`.

```

```

# Print the results

```

```

for result in results:
    print(result)

```

```

{'Neuron Count': 10, 'Mean': 4.468050296490009, 'Standard Deviation':
0.4784714206810787, 'Minimum': 4.279744625091553, 'Maximum':
6.799226760864258, 'Runtime': 26}
{'Neuron Count': 25, 'Mean': 4.373248210320106, 'Standard Deviation':
0.388681963775685, 'Minimum': 4.21902322769165, 'Maximum':

```

```

6.265530586242676, 'Runtime': 26}
{'Neuron Count': 50, 'Mean': 4.335052435214703, 'Standard Deviation':
0.2378184563130353, 'Minimum': 4.202245712280273, 'Maximum':
5.442875862121582, 'Runtime': 26}
{'Neuron Count': 75, 'Mean': 4.390933678700374, 'Standard Deviation':
0.5121579647218228, 'Minimum': 4.183339595794678, 'Maximum':
6.880396842956543, 'Runtime': 26}
{'Neuron Count': 100, 'Mean': 4.338151968442476, 'Standard Deviation':
0.34214244750766304, 'Minimum': 4.188814163208008, 'Maximum':
5.976171016693115, 'Runtime': 26}

```

```

# the optimal number of neurons would be 100.
# This is because the model with 100 neurons in the hidden layer
achieved a lower mean and
# standard deviation of the cost function compared to the other neuron
counts.

```

```

# Make predictions on the test set
y_pred_lstm = model.predict(X_test_reshaped)

# Calculate RMSE
rmse_lstm = np.sqrt(mean_squared_error(y_test, y_pred_lstm))

# Calculate MAE
mae = mean_absolute_error(y_test, y_pred_lstm)

# Calculate R2 score
r2 = r2_score(y_test, y_pred_lstm)

# Print the performance metrics
print("Root Mean Square Error (RMSE):", rmse_lstm)
print("Mean Absolute Error (MAE):", mae)
print("R2 Score:", r2)

```

```

231/231 [=====] - 0s 927us/step
Root Mean Square Error (RMSE): 2.0629785790824524
Mean Absolute Error (MAE): 1.5489273762151918
R2 Score: 0.49857560893330066

```

Compare the performance of the models using RMSE. Which model performed better?

```
# Labels for the models
```

```
models = ['Regression', 'MLPRegressor', 'LSTM']
```

```
# RMSE values
```

```
rmse_values = [rmse_reg, rmse_mlp, rmse_lstm]
```

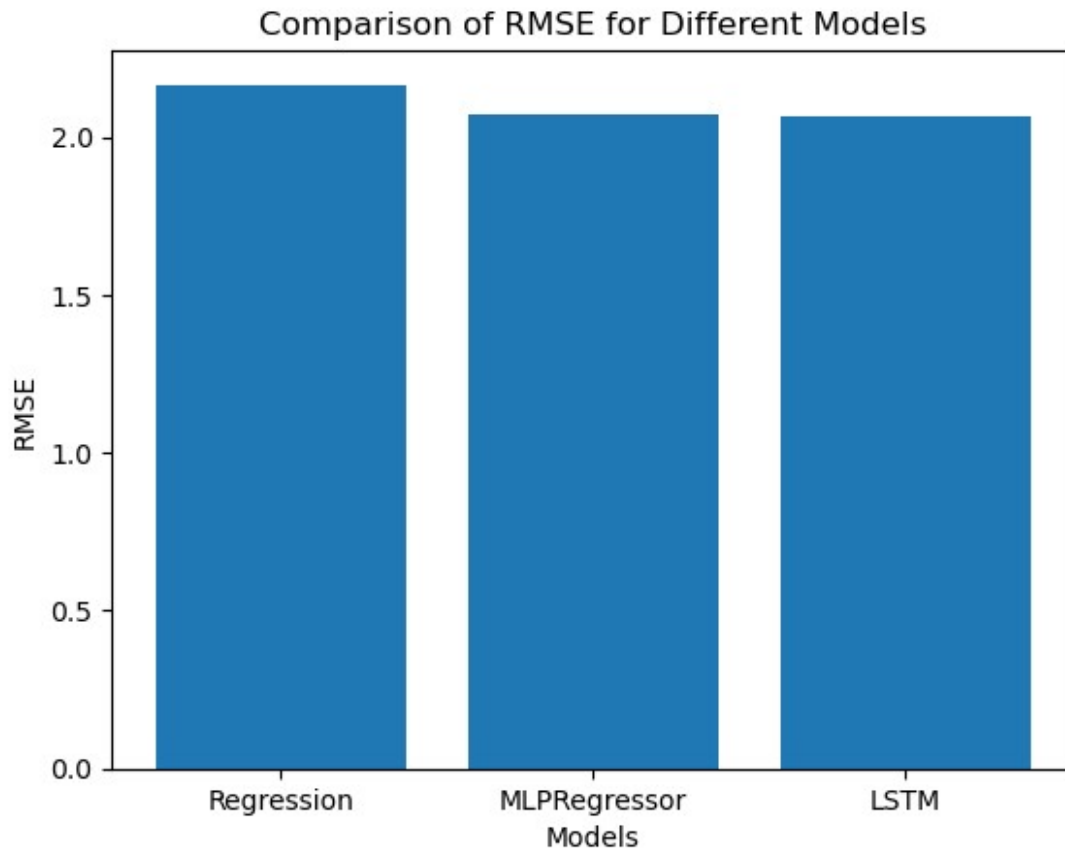
```
# Plotting the bar graph
```

```
plt.bar(models, rmse_values)
```

```
plt.xlabel('Models')
```

```
plt.ylabel('RMSE')
```

```
plt.title('Comparison of RMSE for Different Models')  
plt.show()
```



```
print (rmse_reg)  
print (rmse_mlp)  
print (rmse_lstm)  
  
2.1659689714490153  
2.0713691994393457  
2.0629785790824524
```

###the LSTM model is considered the best model as it has the lowest rmse value